



DAYS 1–5: GENAI CORE

■ DAY 1[part-1] How LLMs REALLY Work (For Developers)

⌚ Goal of PART 1

Before touching any backend or code, **destroy the myth**:

- “LLMs understand code like humans”
- “LLMs predict the next token based on probability”

If this foundation is wrong, **everything later breaks** (RAG, audits, agents).



PART 1 THEORY (CORE MENTAL MODEL)

1 What an LLM ACTUALLY Does (Very Important)

✗ What beginners think

“LLM reads my code, understands logic, finds bugs”

✓ Reality

LLM predicts the next token, again and again

That's it.

No understanding.

No execution.

No compiler.

No debugger.

How generation works (simplified)

When you send this:

```
def add(a, b):  
    return
```

The model internally does:

```
P(token | previous tokens)
```

Example:

- `return` → next token might be:
 - `a`
 - `a +`
 - `a + b`
 - `None`

It picks **one based on probability**.

 This happens **token by token**, not line by line.

Key Truth

LLMs are **probability machines**, not reasoning engines.

Reasoning is an **illusion created by patterns**.

2 Tokens ≠ Words (Why Code Breaks Easily)

✗ Common assumption

“My function is 20 lines, that’s small”

✓ Reality

LLMs don’t see **lines**

They see **tokens**

◊ What is a token?

A token can be:

- Part of a word
- A symbol
- A space
- A newline
- A bracket {
- A tab \t

Example:

```
for(i=0;i<n;i++){
```

This might become:

```
for(|i|=|0|;|i|<|n|;|i|+|+|){
```

⚠ Code = **token-heavy**

Why this matters

- Long files → token explosion
- Minified JS → more tokens
- Nested logic → more context usage

→ This is why “**paste full repo**” fails

3 Why Temperature Matters (Especially for Code)

 Temperature = randomness control

Temperature	Behavior
0.0	Deterministic (same output)
0.2	Safe, boring
0.7	Creative
1.0+	Unstable, risky

For CODE ANALYSIS

- High temperature =  dangerous
- Low temperature =  stable

Why?

Because code analysis needs:

- Consistency
- Accuracy
- Repeatability

Example

At temperature = 0.9:

“This function has SQL injection risk”

At temperature = 0.9 again:

“Looks secure, no issues found”

✗ Same input, different audit → unacceptable

⚡ Why Hallucinations Happen (Critical for CodeAudit)

✗ Hallucination ≠ bug

Hallucination is **expected behavior**

Why hallucinations occur:

1. Model lacks knowledge
2. Context is insufficient
3. Prompt is vague
4. Tokens exhausted
5. Model forced to “continue anyway”

LLMs are **never allowed to say “I don’t know” by default.**

They’d rather:

- ✗ Invent a vulnerability
- ✗ Invent a fix
- ✗ Invent a library

In Code Review

This becomes dangerous:

- Fake vulnerabilities
- Wrong severity
- Incorrect fixes

 This is **WHY RAG exists** (later days)

5 Context Window Limits (Why Long Files Fail)

Context Window = short-term memory

Example:

- GPT-3.5 → ~4k tokens
- GPT-4 → ~8k–32k tokens

What happens when you exceed it?

- Old tokens are dropped
- Important logic disappears
- Model guesses missing context

Common beginner mistake

“I pasted the full backend, why wrong output?”

Because:

- Imports got dropped
- Helper functions vanished

- Security rules missing



MINI CHEAT SHEET (PART 1)

LLMs ≠ thinking machines

LLMs = next-token predictors

Code = token expensive

High temperature = unstable audits

Low temperature = reliable outputs

Hallucination = expected behavior

Context window = hard limit, not suggestion



What you Understood After PART 1

- ✓ Why LLMs feel “smart”
- ✓ Why they fail on long code
- ✓ Why outputs change
- ✓ Why blind trust is dangerous



DAY 1[Part-2] Hands-On: Breaking an LLM (Temperature + Context)



Goal of PART 2

- Same code ≠ same output
- Temperature changes behavior

- Long code causes failure



STEP 0 — Project Structure (IMPORTANT)

Create a fresh folder anywhere:

```
day1-llm-basics/
|
├── main.py
├── requirements.txt
└── README.md    (optional)
```

👉 We are NOT using DB today

Day-1 is **pure LLM behavior**, no distractions.



STEP 1 — Install Dependencies

1. Python -m venv venv
2. Source venv/bin/activate[mac] | venv/scripts/activate[windows]

requirements.txt

Paste this exactly:

```
fastapi
uvicorn
python-dotenv
google-generativeai
```

Now install:

3. pip install -r requirements.txt

STEP 2 — API Key Setup (Critical)

Create a .env file:

```
GEMINI_API_KEY=your_new_rotated_key_here
```

 Never hardcode keys in code.

STEP 3 — First LLM Call (Baseline)

main.py

Paste **everything below :**

```
from fastapi import FastAPI
from pydantic import BaseModel
import google.generativeai as genai
import os
from dotenv import load_dotenv

load_dotenv()

app = FastAPI()

genai.configure(api_key=os.getenv("GEMINI_API_KEY"))

model = genai.GenerativeModel(
    model_name="models/gemini-2.5-flash",
    generation_config={
        "temperature": 0.2
    }
)
```

```
class CodeInput(BaseModel):
    code: str
    temperature: float = 0.2

@app.post("/analyze")
def analyze_code(payload: CodeInput):
    response = model.generate_content(
        contents=[
            "You are a senior backend engineer doing a strict code review.",
            f"Review this code and find issues:\n\n{payload.code}"
        ],
        generation_config={
            "temperature": payload.temperature
        }
    )

    return {
        "analysis": response.text
    }
```

▶ STEP 4 — Run the Server

```
uvicorn main:app --reload
```

Open browser:

<http://127.0.0.1:8000/docs>

👉 You should see Swagger UI.

📝 STEP 5 — Temperature Experiment (IMPORTANT)

Test Code (paste in Swagger):

```
def login(user, password):
    if user == "admin" and password == "1234":
        return True
    return False
```

⌚ Test 1 — Low Temperature

```
{
  "code": "def login(user, password):\n    if user == \"admin\" and password == \"1234\":\n        return True\n    return False",
  "temperature": 0.1
}
```

➡ You'll get:

- Consistent issues
- Same explanation every time

💧 Test 2 — High Temperature

```
{
  "code": "def login(user, password):\n    if user == \"admin\" and password == \"1234\":\n        return True\n    return False",
  "temperature": 0.9
}
```

You'll observe:

- Different vulnerabilities
- Sometimes exaggerated
- Sometimes imaginary

“Nothing changed except **temperature**
Yet audit results changed
This is why AI code audits must be controlled”

Why the SAME code gives DIFFERENT answers

You ran **the same code**, changed **only temperature**:

- Temperature = **0.9**
- Temperature = **0.1**

Yet:

- Both responses were **correct**
- One was **huge, aggressive, creative**
- One was **structured, conservative, predictable**

This is **not randomness by accident**.

This is **how LLMs fundamentally work**.

Core Truth (Drill This)

LLMs do NOT “analyze code”

They predict the NEXT TOKEN based on probability

Everything you saw comes from this.

1 What an LLM is *actually* doing (developer view)

Internally, the model is doing this loop:

```
Given previous tokens →  
Predict probability of next token →  
Sample one token →  
Append →  
Repeat
```

It is **not**:

- running the code
- compiling it
- checking syntax trees
- understanding intent like a human

It's doing:

$P(\text{next_token} \mid \text{all_previous_tokens})$

That's it.

2 Temperature = how much risk the model is allowed to take

💧 Temperature 0.9 (HIGH)

- Probability distribution is **flattened**
- Lower-probability tokens get a chance
- Model becomes:
 - verbose
 - opinionated
 - creative
 - dramatic
 - sometimes *overconfident*

That's why you saw:

- very strong language
- long explanations
- architectural suggestions
- security lectures
- extra code examples

💡 This is dangerous for code review

Because creativity ≠ correctness.

Temperature 0.1 (LOW)

- Probability distribution is **sharp**
- Only high-confidence tokens survive
- Model becomes:
 - conservative
 - repetitive
 - predictable
 - less creative
 - less hallucination-prone

That's why:

- response was calmer
- more structured
- fewer speculative ideas
- safer tone

 **This is preferred for analysis tasks**

3 Why hallucinations happen (VERY IMPORTANT)

Hallucination is NOT a bug.

It's a **mathematical outcome**.

Hallucination happens when:

1. Model is forced to answer
2. Context is incomplete or long
3. Temperature allows exploration
4. The model fills gaps with **statistically plausible tokens**

Example:

- You paste a large codebase
- Important functions fall outside context window
- Model still must predict next token

- It **guesses** patterns it has seen before

➡ Confidently wrong output.

4 Context Window: the silent killer (you already demonstrated it)

When you pasted **short code**, results were good.

If you paste:

- large files
- multiple files
- entire repositories

What happens?

- Older tokens fall out of memory
- Model loses definitions
- Dependencies disappear
- Reasoning collapses

The model **does NOT warn you**.

It will still answer.

This is why:

“Paste entire repo → ask GPT to review”

✗ **fails silently**

5 Why code breaks more easily than text

Tokens ≠ characters ≠ words

Code has:

- symbols
- indentation
- syntax
- long identifiers

So:

- context fills faster
- truncation happens earlier
- one missing function = fake analysis

That's why **LLMs hallucinate more on code than English text.**



Mini Cheat Sheet (Day-1 students MUST remember)

LLMs predict tokens, not truth
Temperature controls creativity vs safety
Low temp = analysis
High temp = brainstorming
Long code breaks context
Confidence ≠ correctness

Q: Which output would you trust more in production?

Ans: Neither blindly — architecture matters more than prompts

This creates **demand for RAG**, which comes next.



How this sets up the rest of the course

Now they understand:

- why naïve GPT usage fails
- why hallucinations exist

- why **systems** matter more than models

Next logical step:

“How do we CONTROL LLMs?”

Which is:

- structured outputs
- chunking
- retrieval
- grounding

★ STEP 6 — “Break the LLM” (Context Limit)

Why this feels confusing → because your brain is thinking like a compiler

But an LLM does NOT work like a compiler

First: Forget everything you know about programs

A compiler / interpreter:

- Reads the **entire file**
- Keeps **all code in memory**
- Errors if anything is missing

✗ LLMs do **NONE** of this

What an LLM ACTUALLY sees

An LLM sees **a sliding window of tokens**, not a file.

That's it.

No AST.

No symbol table.

No project awareness.

What is a Context Window (very important)

Think of context window as:

“How many words the model can see *at one time*”

Example (simplified):

Model context limit = 8,000 tokens

That means:

- It can ONLY “look at” ~8,000 tokens
- If you send more → **older tokens are dropped**

Now let's map this to YOUR demo

You sent this:

```
# utils.py
def helper1(): pass
def helper2(): pass
def helper3(): pass
# repeated 1000+ times
```

Let's say:

- Each function \approx 10 tokens
- 1000 repetitions \approx 10,000+ tokens

 BOOM — context overflow

What ACTUALLY happens inside the model

Step-by-step (this is the key part)

Step 1: Tokens start filling the window

[helper1][helper2][helper3][helper1][helper2] ...

Step 2: Context limit reached

MAX TOKENS REACHED 

Step 3: Old tokens are DROPPED

 First 300 helper functions are gone
 Imports are gone
 Comments are gone

The model **never sees them anymore.**

Critical misunderstanding students have

 “The model read everything but forgot”

NO.

 The model **never sees everything at once**

Why the output looks “dumb” or “wrong”

Now the model sees something like this internally:

```
def helper2(): pass  
def helper3(): pass  
def helper1(): pass  
def helper2(): pass
```

So it thinks:

“Hmm... this looks like repetitive boilerplate. Probably utility functions.”

So it replies with:

- Generic advice
- Fake suggestions
- Hallucinated problems

Why hallucinations happen here

Remember Day-1 rule:

LLMs MUST produce an answer

They are **probability machines**, not truth machines.

So when code is missing:

Unknown logic → Guess likely pattern → Sound confident

That = hallucination.

⚠️ This is NOT a bug (this is crucial)

Why OpenAI / Gemini did NOT “fix” this

Because:

- Context window = memory cost 💰
- Unlimited memory = impossible
- Every model has a hard limit

Even GPT-4, Gemini, Claude — ALL have this.

Real-world consequence (this is why CodeAudit exists)

Naive approach ✗

Paste entire repo → Ask “review my code”

Result:

- Missed bugs
- Fake vulnerabilities
- False confidence

Correct mental model (this is GOLD)

LLM ≠ Code reviewer

LLM ≠ Static analyzer

LLM = Probabilistic text engine

How professionals fix this (preview, not Day-1 yet)

They **never** send full files.

They:

- Chunk code
- Summarize chunks
- Build RAG
- Track context manually

👉 That's literally what your **CodeAudit project** is about.

One-line explanation (memorize this)

"If the code doesn't fit in context, the model guesses — and guessing looks like hallucination."



MINI CHEAT SHEET (PART 2)

Temperature controls randomness

High temp = creative + unstable

Low temp = boring + reliable

Long code ≠ fully read code

Context window is a HARD LIMIT

LLMs guess when unsure

They don't say "I don't know"

Long input → context overflow

Context overflow → token dropping

Token dropping → missing awareness

Missing awareness → hallucination

After PART 2, You will be aware off

- ✓ Why same code gives different audits
- ✓ Why production GenAI must control temperature
- ✓ Why full-repo paste is useless
- ✓ Why CodeAudit NEEDS chunking + RAG