

Part 2 was planned as:

```
| .env.development    ← settings for local development
| .env.production     ← settings for production/cloud deployment
| Config validation    ← check if all required env vars exist on startup
```

But here's the thing — we've already done most of this:

```
✓ .env                → your local development settings (already exists)
✓ .env.docker         → your Docker/production settings (just created)
✓ Config validation → Pydantic Settings in config.py already validates
                    (if GEMINI_API_KEY is missing, app crashes with clear
                    error)
```

So Part 2 is essentially **already complete** from our Docker setup. We don't need to create separate files just for the sake of it.

Let's move to **Part 3: Logging & Monitoring** — that's where real value is. It adds structured logging (instead of random `print()` statements), health check endpoints, and error tracking. This is something interviewers love to see.